

MPP coursework

Details and answers to common questions

Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

http://creativecommons.org/licenses/by-nc-sa/4.0/deed.en_US

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Acknowledge EPCC as follows: “© EPCC, The University of Edinburgh, www.epcc.ed.ac.uk”

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.

Edge Detection

- Represent pixels in an $M \times N$ image as a 2D array of values $image_{i,j}$ in the range 0 (black) to 255 (white)
 - i is in the range 1, 2, ..., M ; j is in the range 1, 2, ..., N
 - simple to store as an ASCII PGM file
- Apply a simple edge detection algorithm – combine each pixel with its 4 nearest neighbours

$$edge_{i,j} = image_{i+1,j} + image_{i-1,j} + image_{i,j+1} + image_{i,j-1} - 4 image_{i,j}$$

- What about the pixels at the edges since neighbours are outside the original image?
 - set the boundary conditions so all these values are 255 (white)

Image reconstruction

- We choose to reconstruct the image from the edges
 - turns out that this *is* possible given the way we compute the edges
- Algorithm
 - make an initial guess at the image values old_{ij}
 - update *many times* using this equation:

$$new_{i,j} = \frac{1}{4} (old_{i+1,j} + old_{i,j+1} + old_{i-1,j} + old_{i,j-1} - edge_{i,j})$$


- copy *new* back to *old*
- repeat ...

Edge detection / image reconstruction



single
pass

hundreds of
iterations



Serial code



- Uses explicit halos for boundary conditions
 - must use the same boundary conditions used when creating edges
 - declares arrays to have bounds $0, 1, \dots, M+1$ and $0, 1, \dots, N+1$
 - fills boundaries (all pixels with $i=0$, $i=M+1$, $j=0$ or $j=N+1$) with 255
 - reads the input (the edge) into an array *buf* with no halos
 - copies *buf* to the centre of the *edge* array
 - applies the equation many times (1500 iterations)
 - copies *old* back to *buf* excluding the halos
 - writes *buf* to an output file
- Simplifications
 - image dimensions and file names specified at compile time

Parallel algorithm

- Split image into slices between all processes
 - divide across the first index i in C, second index j in Fortran
 - local array sizes are now MP and NP (plus halos)
 - C: $MP = M/P$, $NP = N$; Fortran: $MP = M$, $NP = N/P$
- Rank 0 reads into $M \times N$ *masterbuf* array (with no halos)
 - scatter to local *buf* array
- Implement exactly the same calculation as before except:
 - loop limits are MP and NP (not M and N)
 - initialise all halo values to 255
 - **before start** of each iteration, swap boundary values as appropriate
- Gather *buf* array back to *masterbuf*
 - Rank 0 writes to file

Domain Decomposition

- Different choices in C and Fortran



Parallel Case Study solution

- Simplifications
 - specifies number of processes P at compile time
 - to make array allocation simpler
 - assumes that the image divides exactly between P
 - doesn't even check that this is true!
 - uses MPI_Sendrecv
 - you should use **explicit** non-blocking, i.e. MPI_Isend and/or MPI_Irecv

Coursework

- Differences from Case Study
 - use a 2D decomposition, i.e. try to split into many rectangular regions, not restricted to slices
 - requires more halo swaps than before (and they must be non-blocking!)
 - you cannot do the IO using scatter and gather
 - different boundary conditions
 - fixed values for $i = 0$ and $i = M+1$ (as before), but now a function of i
 - periodic for $j = 0$ and $j = N+1$: $image_{i,N+1} = image_{i,1}$; $image_{i,0} = image_{i,N}$
 - have therefore supplied new edge files, e.g. edgenew192x256.pgm
 - new reference serial codes: imagenew.c, imagenew.f90
 - compute the average pixel value at regular intervals
 - stop when Δ parameter is smaller than some threshold (e.g. 0.1)
 - equals the maximum absolute change in any pixel between *new* and *old*

Questions

- What if the image does not decompose exactly?
 - you might want to deal with this case (see tutorial problems)
 - however, perfectly acceptable to quit and say “can’t do this!”
 - can still run on many different process counts
- What about a prime number of processes?
 - on 3 processes, for example, your code should just use a 1x3 or 3x1 decomposition (i.e. slices as for the case study)
 - 1D decomp is a subset of 2D: your code should run without problems
- Should I get the same answer?
 - Yes! The exercise has been designed so that (rather unusually for a parallel program) I would expect you to get exactly the same answer regardless of how many processes you run on

Hints (i)

- Do no worry about doing elegant IO
 - just implement something simple that works
 - e.g. broadcast the whole image to all processes
 - copy relevant section to your local edge array
- Test section by section
 - run on a single process
 - checks for serious bugs!
 - run for zero iterations
 - checks the IO is correct
 - run with no halo swaps
 - checks the calculation is correct
 - run with only horizontal or vertical halo swaps

Hints (ii)

- You don't need to run for thousands of iterations
 - can set a maximum count for performance testing
 - or use a larger value for Δ
- You are welcome to use your own (larger) images
 - I will upload instructions
- Bring your problems to the tutorial and drop-in sessions!